

The background of the slide is a light gray gradient, decorated with numerous realistic water droplets of various sizes. Some droplets are at the top, some at the bottom, and some on the right side, creating a clean, modern aesthetic.

COLUMN-STORES VS. ROW-STORES: HOW DIFFERENT ARE THEY REALLY?

DANIEL J. ABADI (YALE)

SAMUEL R. MADDEN (MIT)

NABIL HACHEM (AVANTGARDE)

PRESENTATION BY PRANAV GOEL

Introduction

- On analytical workloads, Column store are found to perform order of magnitude better than traditional row-oriented database systems – “row stores”
- Elevator pitch: “column-stores are more I/O efficient for read-only queries since they only have to read from disk (or from memory) those attributes accessed by a query”
- Factor of 2 on price/performance
- Factor of 5 on performance



Data Warehouse and DBMS Software

- \$5 Billion Data warehouse industry out of 20 Billion DBMS software industry of recurring revenues in 2010
- Growing at 8.5 % annually
- Total industry is expected to be at 105 Billion by 2020

Source Gartner



Motivation

*Why not traditional DBMS
system move to C-Store?*

C-Store:

SAP HANA

Sybase IQ

Vertica




Key Questions:

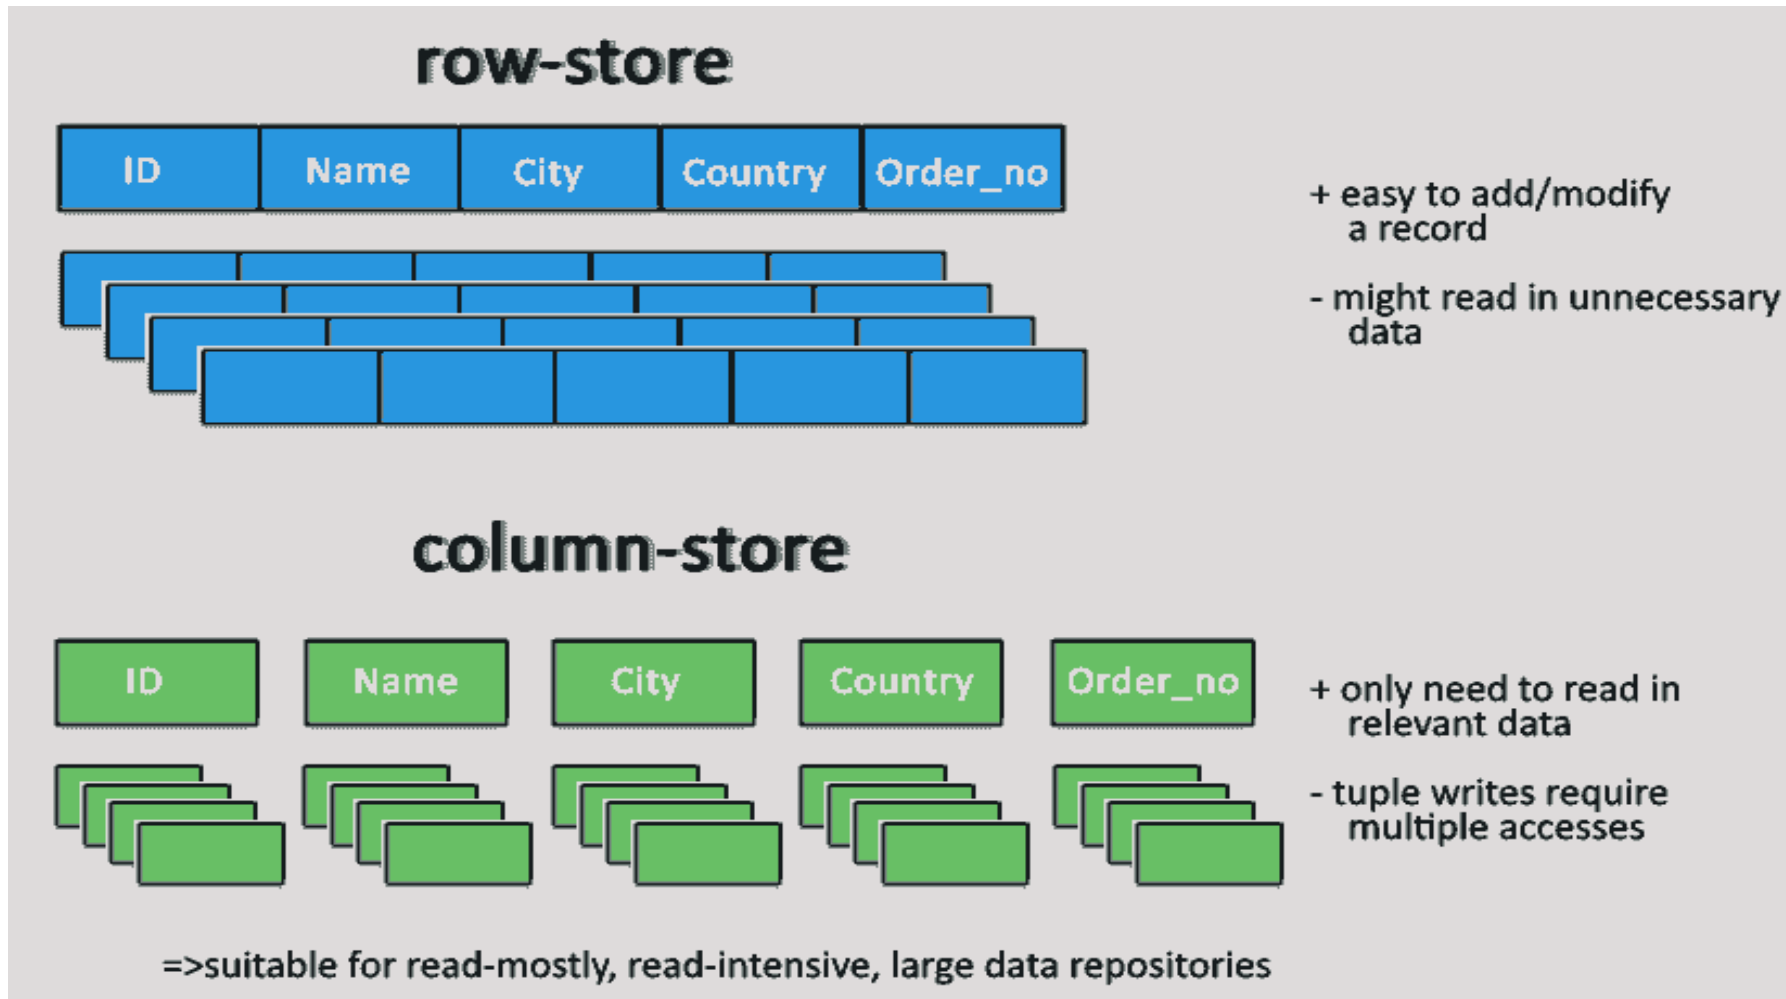
- How much of the buzz around column- stores just marketing hype?
- Do you really need to buy products like SAP HANA or Vertica?
- How far will your current row-store take you?

Is this assumption valid?

- Can we adapt our row-store to get column-store performance?
- Can you simulate a column-store in a row-store?
- If not, what makes column-store so special?



Background



Background: Continued

Most of the queries do not process all the attributes of a particular relation.

- For example the query

Select c.name and c.address

From CUSTOMER as c Where c.region = 'Vancouver';

- Only processes three attributes of the relation CUSTOMER. But the customer relation can have more than three attributes.
- Column-stores are faster for OLAP operations and slower for OLTP operations with many row inserts

Paper Methodology

- Comparing row-store vs. column-store is dangerous/borderline meaningless
- Instead, compare row-store vs. row-store and column-store vs. column-store
 - ✓ Simulate a column-store inside of a row-store
 - ✓ Remove column-oriented features from column-store until it behaves like a row-store

Row Oriented Execution

- Common assumption would be to based on the storage layout, is to modify the row store physical structure in such a way that it behaves more like column store :
- The three ways to do these are:
 - ✓ Vertical Partitioning
 - ✓ Using index-only plans
 - ✓ Using materialized views


First: Vertical Partitioning

- Idea: Vertical partition every relation
 - Mechanism to connect files from same row together :
 - This is done by addition **integer** 'position' column to every table
 - Primary keys are bad idea as they are could be long or **composite**
 - Thus each column have one Physical table.
 - Join are done based on position attributes when fetching multiple columns

Problems:

- “Position” - Space and disk bandwidth
- Header for every tuple – further space wastage

Vertical Partitioning : Example

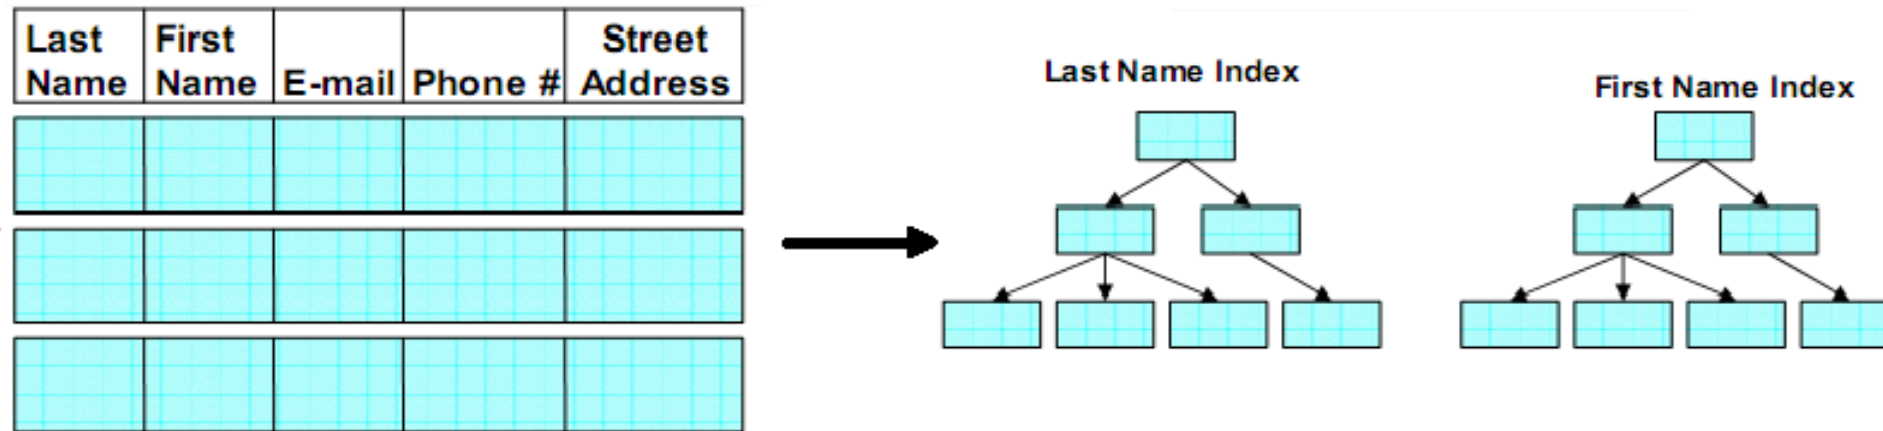


Vertical Partitioning

Second: Index-Only Plans

- Approach
 - Create B+ tree like index for every column
 - Plans never access the actual tuples on disk
 - Tuple headers are not stored thus less overhead

Example:



Index Every Column

Index-Only Plans

Problems

- Separate indices may require full index scan, which is slower

Example:

```
SELECT AVG(salary) FROM emp WHERE age > 40
```

As there are separate indices on age and salary, an index only plan have to find ids corresponding to satisfying age and then merge with (id,salary) extracted from salary index

- Slow tuple construction

Third: Materialized Views(MV)

Methodology:

- Create an optimal set of materialized views for every query workload.

Idea behind Strategy:

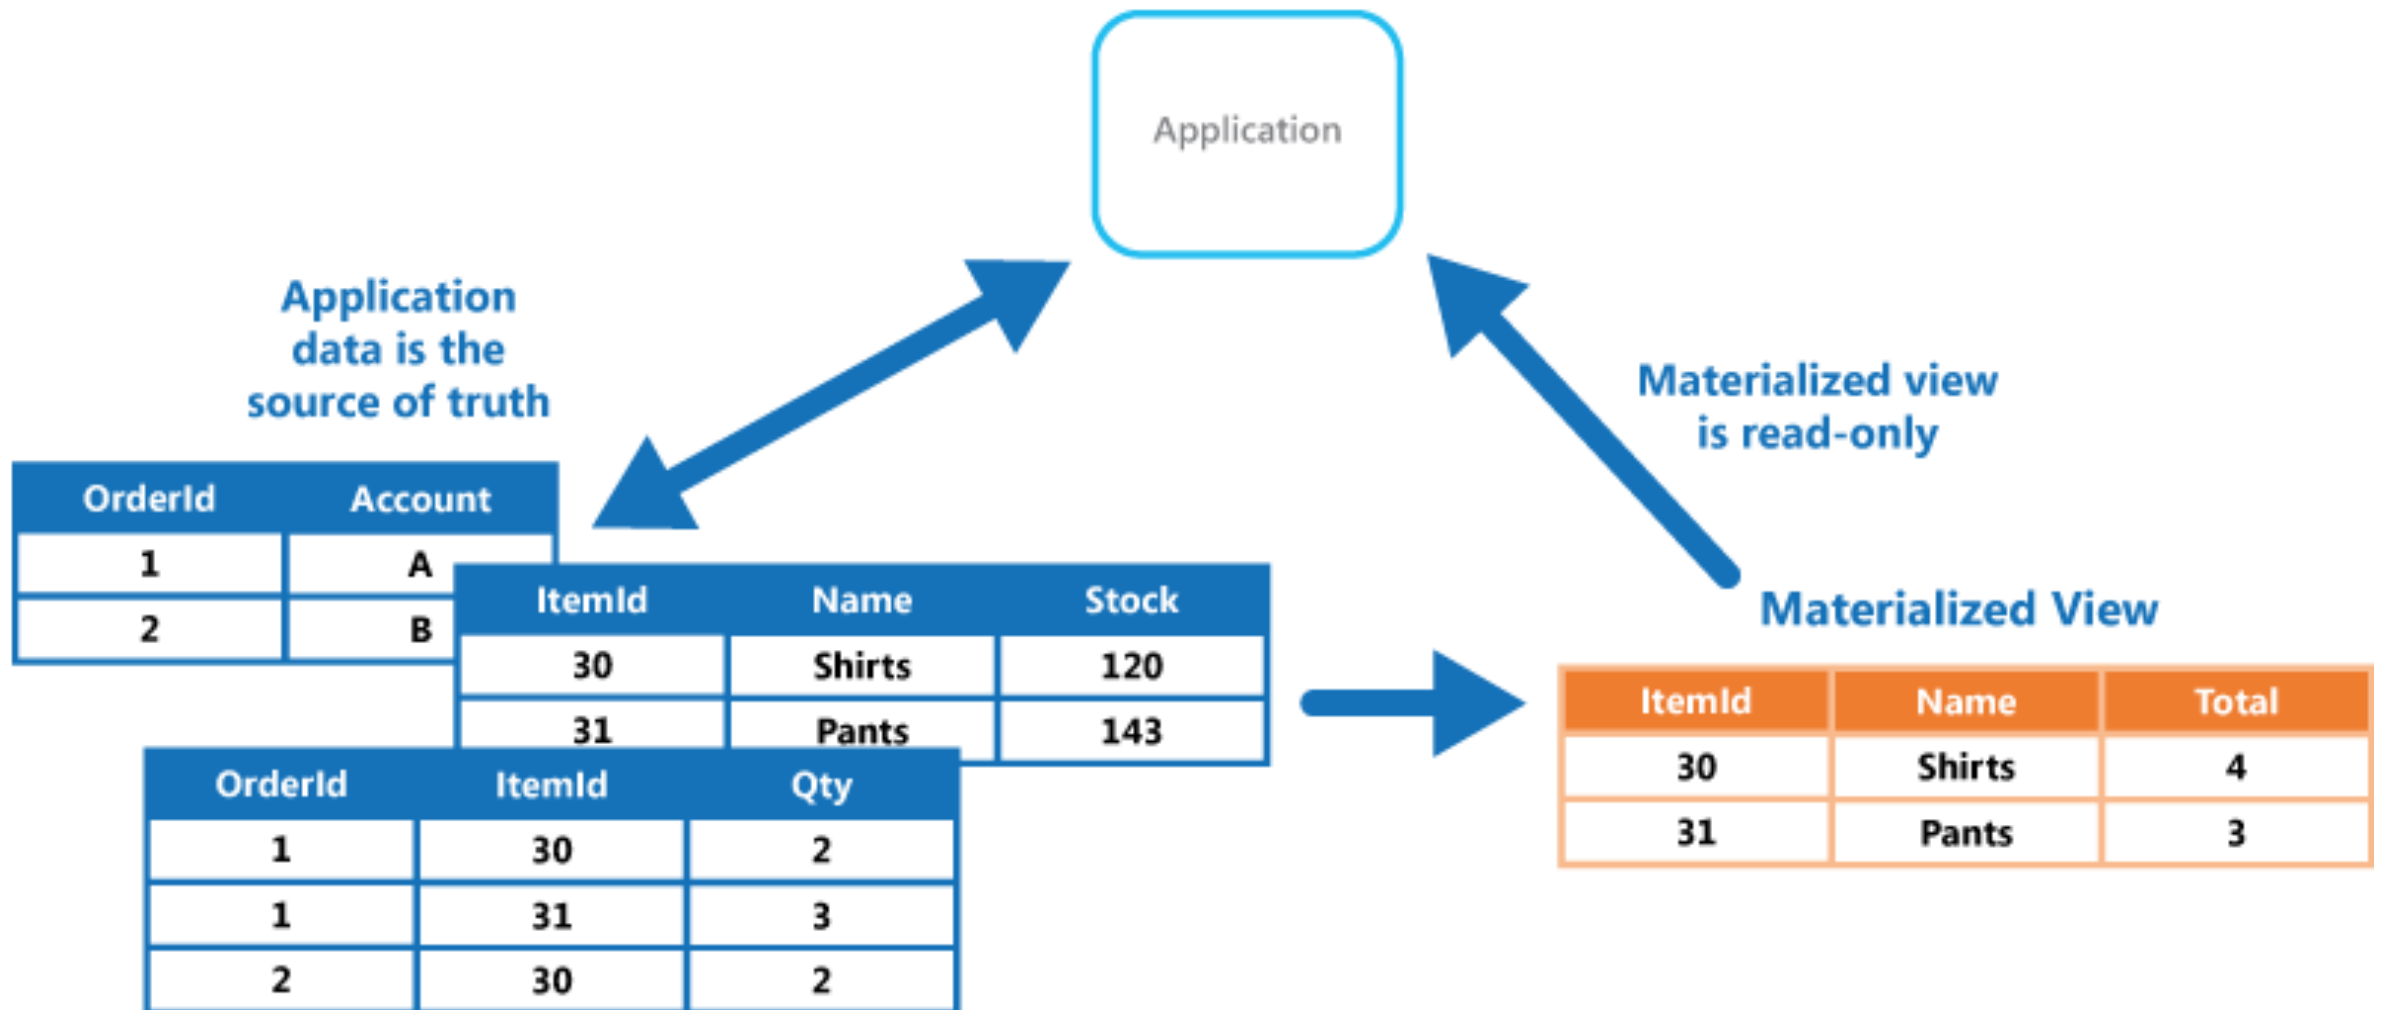
Only have columns needed for answer

Avoid overhead – perform better

Problems:

Advance knowledge of query workload thus practical in limited situations

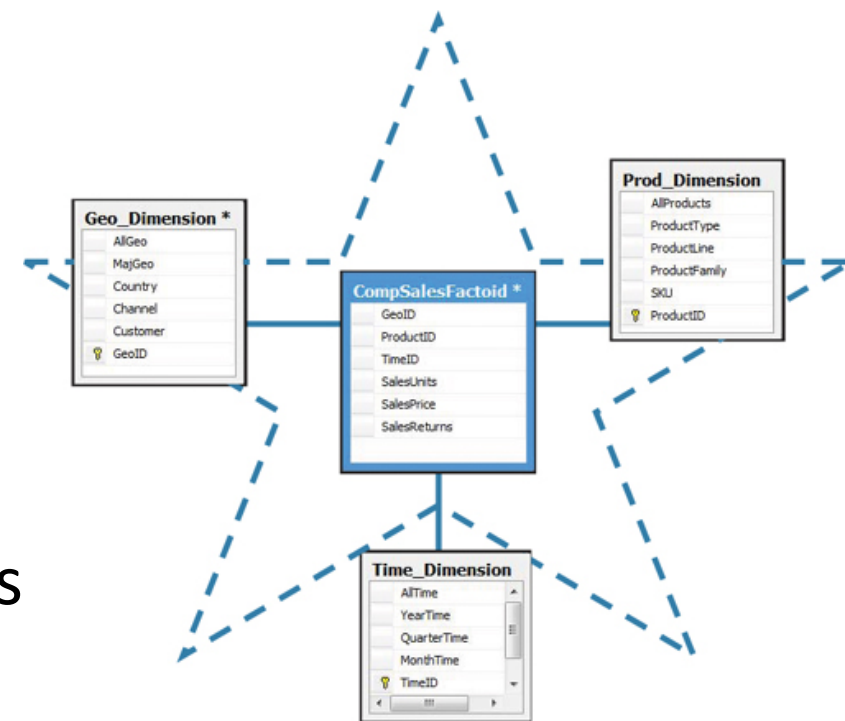
Materialized Views: Example



Benchmarking Methodology

Star Schema Benchmark (SSBM)

- Fact table contains 17 columns and 60,000,000 rows
- 4 dimension tables, biggest one has 80,000 rows
- Queries perform 2-4 joins between fact table and dimension tables, aggregate 1-2 columns from fact table
- All benchmarks are done on using similar hardware configuration named **SYSTEM X** for this paper




Performance Comparison Row Store

Average case performance across all workloads of the SSBM

T – Traditional
T(B) - Traditional Bitmap
MV – Materialized Views
VP – Vertical portioning
AI – All indexes

- MV performs best
- Index only are the worst



Column Store simulation in Row Store: Analysis

Vertical Partitioning:

- **Tuple overheads:**

- LineOrder Table

60 million tuples, 17 columns

- 8 bytes of over head per row

- 4 bytes of record-id

Tuple Header	TID	Column Data
xxxx	1	yyy
xxxx	2	yyy
xxx	3	yyy

	1 column (Compressed)	Whole table (Compressed)
Row Store	0.7 -1.1 GB	4 GB
Column Store	240 MB	2.3 GB

Column Store simulation in Row Store: Analysis

- All indexes approach is a poor way to simulate a column-store
- Problems with partitioning are fundamental – Store tuple header in a separate partition
 - Allow virtual TIDs
 - Combine clustered indexes, vertical partitioning
- **What can possible be done to simulate column store in row store?**
 - Need better support for partitioning at the storage layer
 - Need support for column-specific optimization at the executor level
 - Full integration: buffer pool, transaction manager

Column-Oriented Execution

- Different optimization for column oriented database –
 - ✓ Compression
 - ✓ Late Materialization
 - ✓ Block Iteration

The **idea** is to minimize these advantages of column store to simulate row store in column store

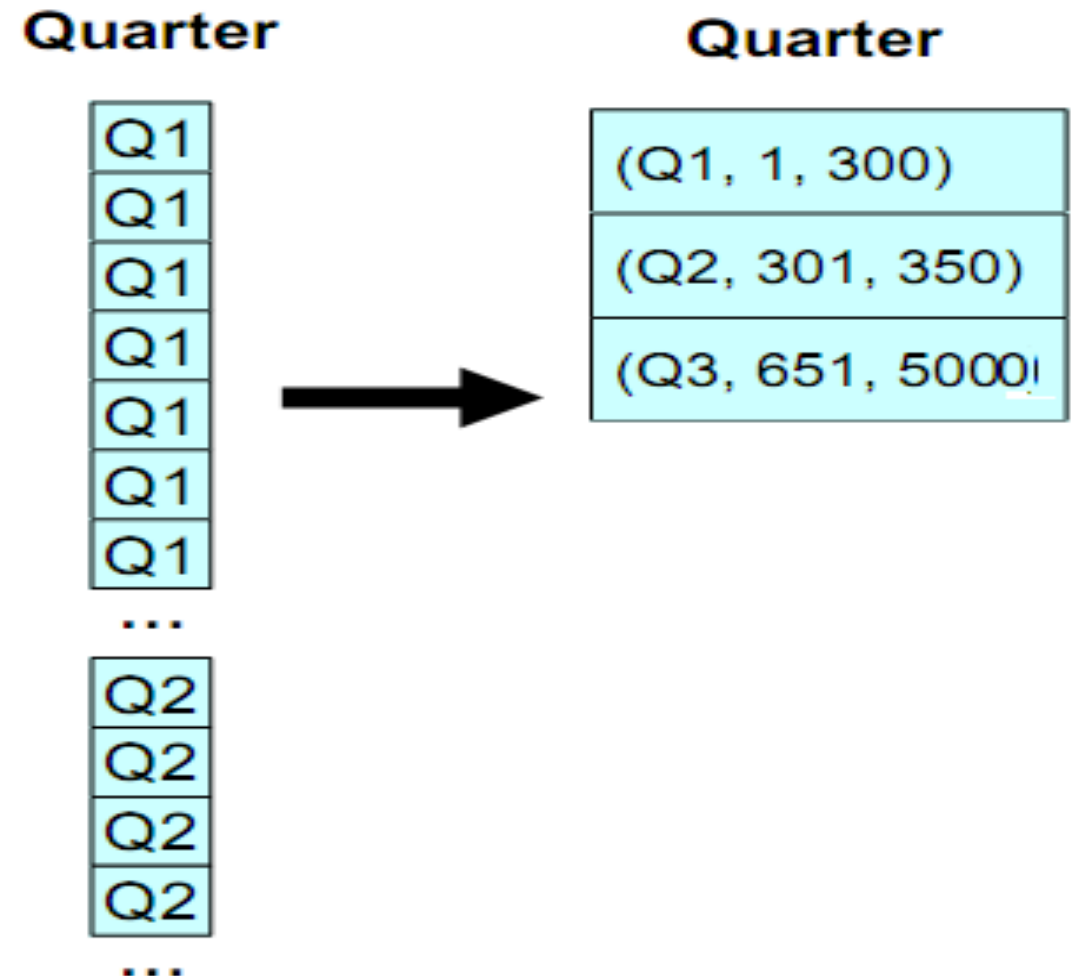
Compression:

- Low information entropy (high data value locality) leads to high compression ratio

- Advantages

- Disk Space is saved
- Less I/O
- CPU cost decrease if we can perform operation without decompressing

Idea: Use unsorted data



Late Materialization:

- Most queries require data from multiple columns that be combined together into rows to form information about an entity.
- **Tuple Construction:** Join like materialization of tuples
- Delay tuple construction

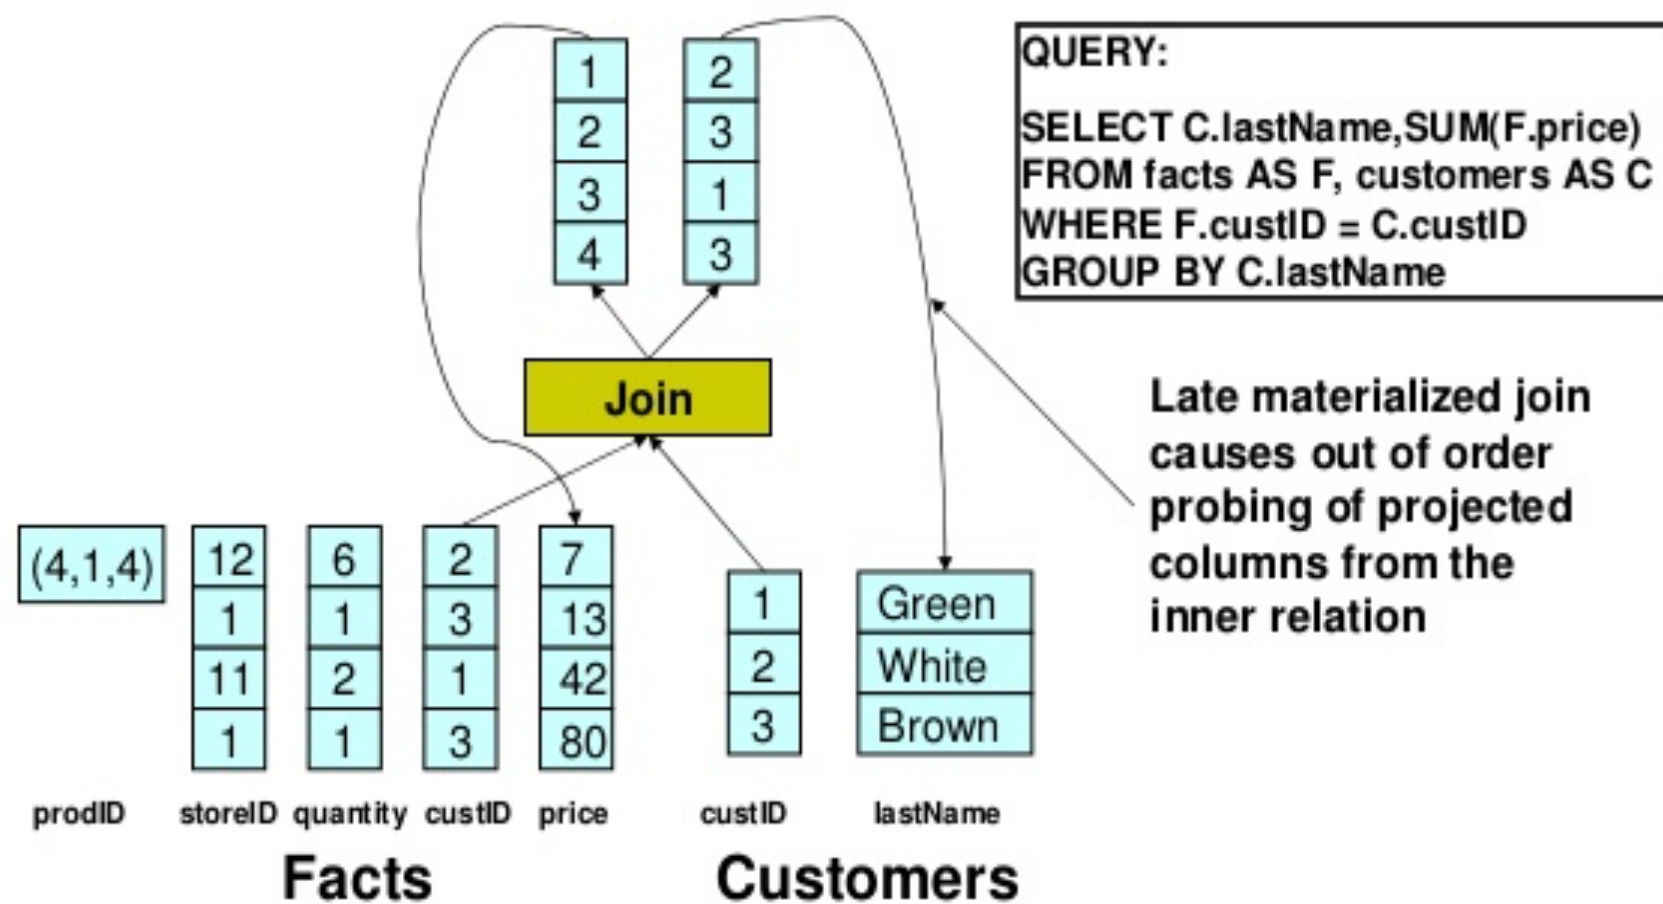
Advantages

- Unnecessary Construction of tuple is avoided
- Direct operations on compressed data
- Cache performance is improved

Idea: Use Early Materialization in the query plan



Late Materialization Example



Block Iteration

Row Store first iterate through each tuple to extract needed attributes thus leading to **tuple at a time processing** (MySQL).

C-Store blocks of values from same column are sent to operator in single function call

- Iterate over blocks rather than tuples, like Batch Processing

Advantages:


- Minimizes per-tuple overhead – Exploits potential for parallelism

Idea : Can be applied even in Row stores – IBM DB2 implements it

Experiment:

Performance of C store with various optimizations removed

T=tuple-at-a-time processing;
t=block processing;
I=invisible join enabled;
i=disabled;
C=compression enabled,
c=disabled;
L=late materialization enabled;
l=disabled;



CONCLUSION:

- TO SIMULATE COLUMN STORE IN ROW STORE, TECHNIQUES LIKE –
VERTICAL PARTITIONING , INDEX ONLY PLAN

DO NOT YIELD GOOD PERFORMANCE

HIGH PER-TUPLE OVERHEADS, HIGH TUPLE CONSTRUCTION COSTS ARE THE REASON

- WHERE AS IN COLUMN STORE –
LATE MATERIALIZING , COMPRESSION , BLOCK ITERATION

ARE THE REASONS FOR GOOD PERFORMANCE

MIGHT BE POSSIBLE TO SIMULATE A ROW-STORE IN A COLUMN-STORE,

NEED BETTER SUPPORT FOR VERTICAL PARTITIONING AT THE STORAGE LAYER

NEED SUPPORT FOR COLUMN-SPECIFIC OPTIMIZATIONS AT THE EXECUTER LEVEL

- DO WE SEE COLUMN STORE SIMULATION IN ROW STORE WORKING?
- WILL COLUMN STORE PROVIDE WRITE OPTIMIZATIONS COMPARABLE TO ROW STORES?
- USE ROW STORE AND COLUMN STORE INTERCHANGEABLY

SAP HANA – USES CONCEPT OF DELTA MERGE

WHERE WRITES ARE DONE ON ROW STORE AND THEN MERGED WITH COLUMNS STORE





THANK YOU FOR
YOUR LISTENING

DO YOU HAVE
ANY QUESTIONS?