

Shark: SQL and Rich Analytics at Scale

Guiliang Liu
2017/02/13

Outline

- Abstract
- Motivation
- Introduction
- System Overview and Engine Extension
- Experiment
- Discussion about shark

Abstract

- Built on the Spark
- Compatible with Hive
- A MapReduce-like execution engine
- Run SQL queries and sophisticated analytics functions



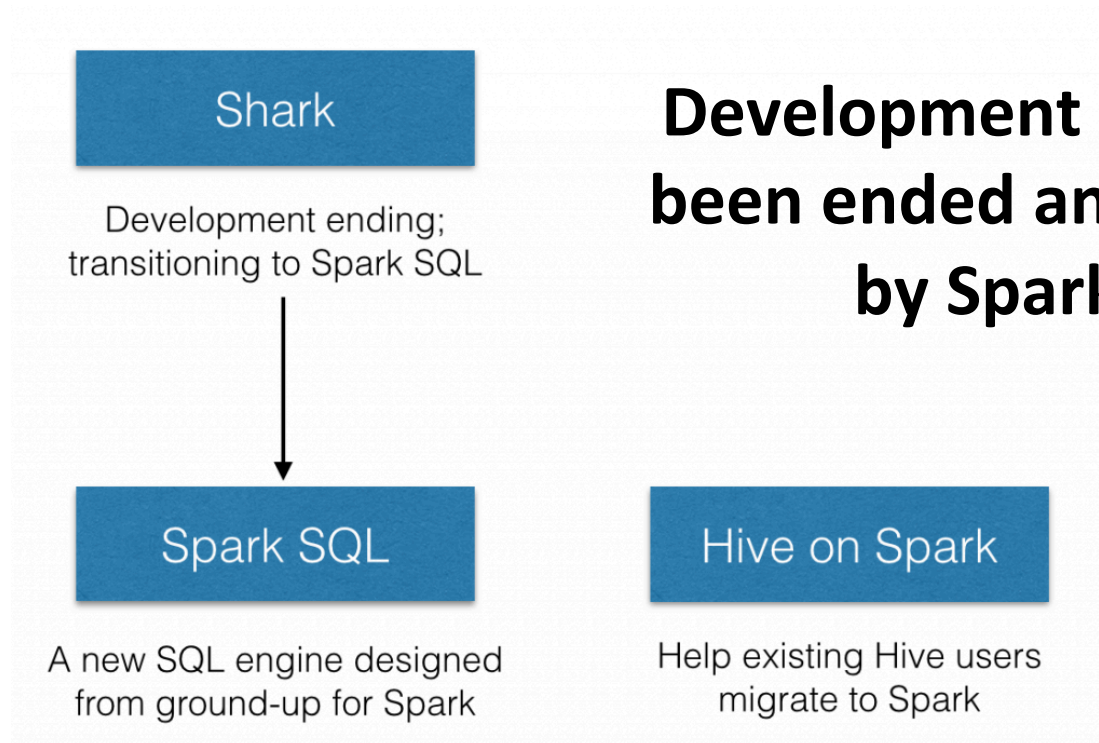
Abstract

- Run SQL queries up to 100× faster than Apache Hive
- Run machine learning programs more than 100× faster than Hadoop.



However

- When you go to Sharks project in AMP LAB



Why

- Why Shark has been abandoned ?
- What's the influence of Shark in the development of modern database?
- What did Shark give us?



Motivation

- To tackle the “big data” problem

1. MapReduce

- (1) A fine-grained fault tolerance model

- (2) Latencies is huge

- (3) Try to optimized MapReduce for SQL queries.

2. MPP analytic databases

- (1) Employ a coarser-grained recovery model

- (2) Works well for short queries where a retry
is inexpensive

- (3) Faces significant challenges for long queries



Motivation

- Why do we develop Shark?
 1. support both SQL and complex analytics efficiently
 2. provide fine-grained fault recovery across both types of operations.
 3. Compatible with popular system like Hive



Introduction

- To be specific

1. Built on Resilient Distributed Datasets (RDDs)

- (1) Perform most computations in memory
- (2) Offering fine-grained fault tolerance

2. Built on Spark and added several features

- (1) In-memory columnar storage and columnar compression
- (2) Advance Spark with Partial DAG Execution (dynamic mid-query replanning)

3. Compatible with Apache Hive

- (1) Supporting all of Hive's SQL dialect and UDFs
- (2) Allowing execution over unmodified Hive data warehouses

Introduction

- Significance of Spark

1. **Speed**

Shark can answer SQL queries and run **iterative** machine learning algorithms faster

2. **Exploration**

(1) Shows that MapReduce-like execution models can be applied effectively to SQL (Debate).

(2) Offer a promising way to combine relational and complex analytics.

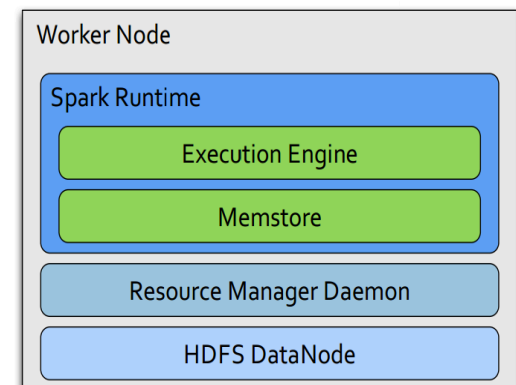
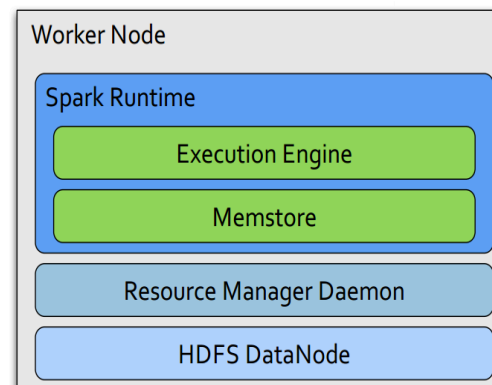
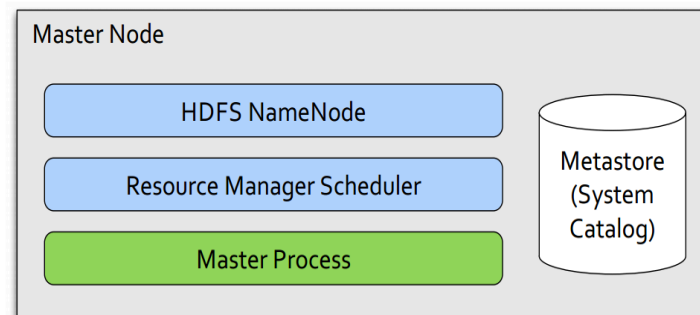
System Overview

1. Compatible with Hive

- (1) Support Hadoop storage API (Like HDFS, Amazon EC2)
- (2) Support a wide range of data formats (like Json, XML, binary sequence file)

2. Built on top of spark

- (1) Compile query into operator tree represented as RDD
- (2) Cluster resources can optionally be allocated by a resource manager (Hadoop Yarn, Apache Mesos)



System Overview

3. Fault Tolerance provided by RDD

- (1) Shark remembers the lineage of the RDD
- (2) Shark can tolerate the loss of any set of worker nodes
- (3) Recovery is parallelized across the cluster
- (4) Recovery works even for machine learning UDFs


4. Executing SQL over RDD

- (1) The same as Hive in query parsing, logical plan generation

- (2) Different from Hive in physical plan generation

Shark:

transformations on RDDs rather than MapReduce jobs



Engine Extension

1. Partial DAG Execution (PDE)

Motivation:

1. Query fresh data that has **not** undergone a data loading process
2. **Precludes** static query optimization techniques

Originally:

1. Spark materializes the output of each map task in memory before a shuffle,
2. Only spilling it to disk as necessary
3. Reduce tasks fetch this output.

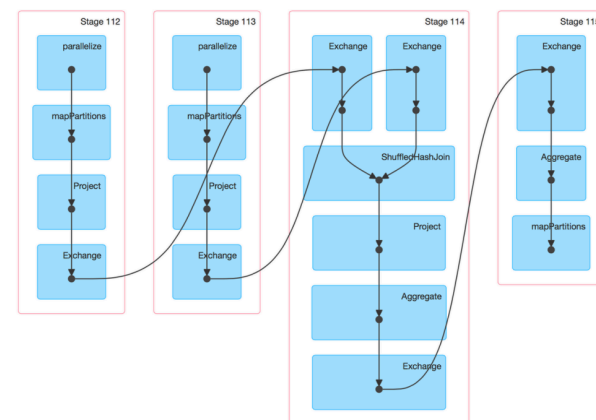
Modification:

1. While materializing map outputs, it gathers customizable statistics at global and per-partition granularities
2. it allows the DAG to be altered based on these statistics (aggregated and presented)

Details for Job 8

Status: SUCCEEDED
Completed Stages: 4

► Event Timeline
▼ DAG Visualization



Engine Extension


1. Partial DAG Execution (PDE)

Application Example: Join Optimization

We have two kinds of join:


(a)map join

- A small input table is broadcast to all nodes
- Joined with each partition of a large table
- Avoid repartitioning and shuffling phase



(b)shuffle join

- Hash-partitioned by the join key
 - reducer joins corresponding partitions
- using a local join algorithm




Engine Extension

1. Partial DAG Execution (PDE)

Application Example: Join Optimization

How to decide which join to apply?

- Map join can avoid repartitioning and shuffling phase by broadcasting table.
- But if the table is too large, very inefficient
- Map join is only worthwhile if some join inputs are small



Shark uses partial DAG execution to select the join strategy at runtime based on its inputs' exact sizes.

- **Map Join:** If one of the tables is rather small compared to other table
- **Shuffle Join:** both of the tables are really large

Engine Extension

2. Columnar Memory Store

In-memory data representation is important for in-memory computation

(a) Spark: store data partitions as collections of JVM objects (Large Overhead)

(b) Shark: store all columns of primitive types as JVM primitive arrays.

(Similar to columnar database systems, e.g., Cstore)

3. Data Co-partitioning

(a) In some warehouse workloads, *two tables are frequently joined together*

(b) co-partition the two tables based on their *join key* in the data loading process

(c) Put them in the same node

(d) Reduce data transferring time

(main consumption in distribution system)



Machine Learning Support

Shark supports machine learning as a first-class citizen.

- Choose Spark as the execution engine
- Choose RDD as the main data structure for operators.

Some features (Two integrations):

(1) Language Integration

- In addition to SQL query, Shark also allows queries to return the RDD representing the query plan.
- Shark can then invoke distributed computation over the query result using the returned RDD.
- Example: **logistic regression**

(2) Execution Engine Integration

- machine learning computations and SQL queries to share workers and cached data
- Data transferred conveniently (Accelerate)

```
def logRegress(points: RDD[Point]): Vector {  
  var w = Vector(D, _ => 2 * rand.nextDouble - 1)  
  for (i <- 1 to ITERATIONS) {  
    val gradient = points.map { p =>  
      val denom = 1 + exp(-p.y * (w dot p.x))  
      (1 / denom - 1) * p.y * p.x  
    }.reduce(_ + _)  
    w -= gradient  
  }  
  w  
}  
  
val users = sql2rdd("SELECT * FROM user u  
JOIN comment c ON c.uid=u.uid")  
  
val features = users.mapRows { row =>  
  new Vector(extractFeature1(row.getInt("age")),  
    extractFeature2(row.getStr("country")),  
    ...)  
}  
val trainedVector = logRegress(features.cache())
```

Experiments

Claims:

Sharks perform more than **100× faster** than Hive and Hadoop

Sharks thought himself very strong
Let's use experiment to prove it!



Machine Parameters:

(1) experiments were conducted on Amazon EC2 using 100 m2.4xlarge nodes.

Each node had 8 virtual cores, 68 GB of memory, and 1.6 TB of local storage.

(2) The cluster was running 64-bit Linux 3.2.28, Apache Hadoop 0.20.205, and Apache Hive 0.9.

Experiments

1. Pavlo et al. Benchmarks

The benchmark used two tables:

- (1) 1 GB/node rankings table,
- (2) 20 GB/node user-visits table.

pay attention to Memory-based Shuffle

Compare performance of

- (1) Shark (2) Shark(disk) (3) Hive
- by measuring the time to load data into HDFS and Shark's Memory store

Experiment steps

- (1) Selection Query (2) Aggregation Quer
- (3) Join Query (4) Data Loading

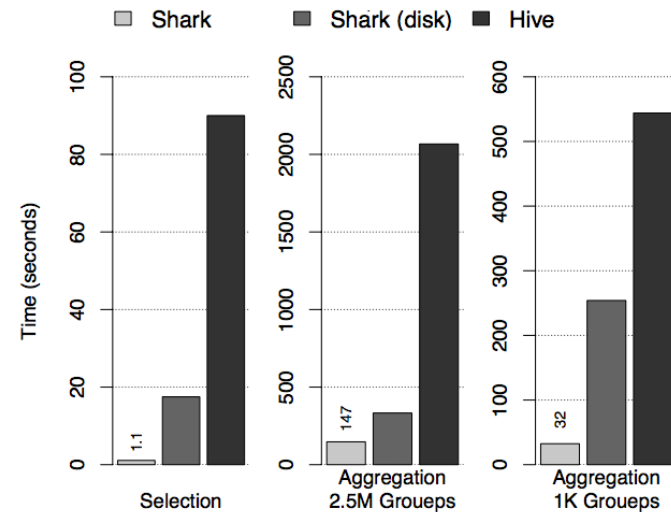


Figure 4: Selection and aggregation query runtimes (seconds) from Pavlo et al. benchmark




Figure 5: Join query runtime (seconds) from Pavlo benchmark

Experiments

2. Micro-Benchmarks

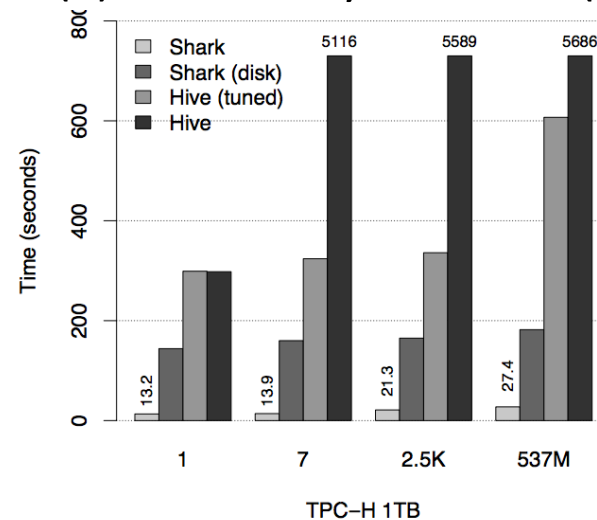
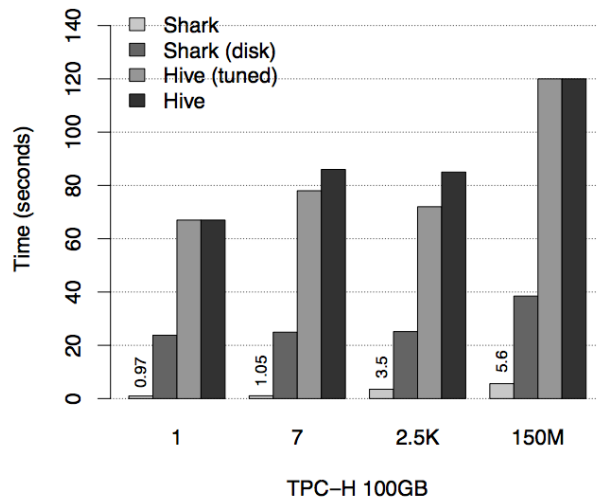
The Micro-benchmark contains:

100 GB and 1 TB of data generated by the DBGEN program provided by TPC-H

Experiment steps:

(1) Aggregation Performance


To measuring Shark's performance on both (a) in-memory data and (b) data loaded



Experiments


(2) Join Selection at Run-time

In this experiment, we tested how **partial DAG execution** can improve query performance through run-time re-optimization of query plans.



(3) Fault Tolerance

- To measure Shark's performance in the presence of node failures.
- measured query performance before, during, and after failure recovery.



Experiments

3. Real Hive Warehouse Queries

- To test the performance of sharks in the real industry environment.
- Use a sample of their Hive warehouse data with two years of query traces from Hive system.

4. Machine Learning

Implemented two machine learning algorithms
(1) logistic regression (2) k-means

To compare the performance of Shark versus running the same workflow in Hive and Hadoop.

Three steps:

- (1) Select data from warehouse using SQL
- (2) Extracting features (3) Applying algorithm


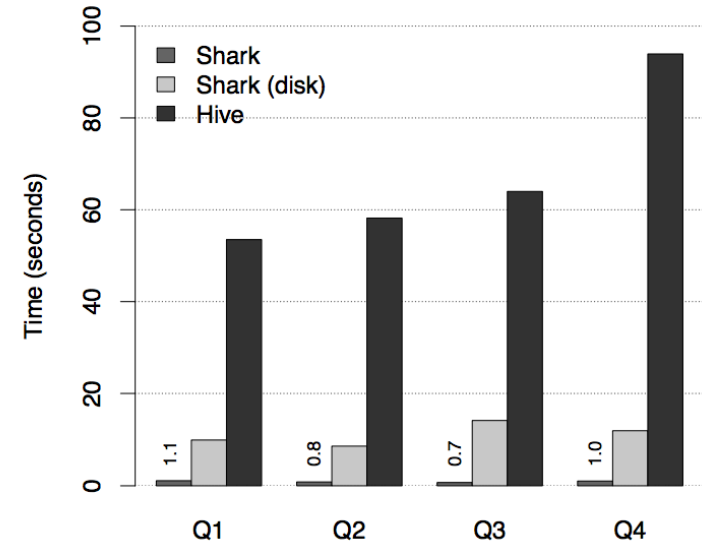


Figure 10: Logistic regression, per-iteration runtime (seconds)




Figure 11: K-means clustering, per-iteration runtime (seconds)

Discussion

Why are Previous MapReduce-Based Systems Slow?

Intermediate Outputs:

- In MapReduce, map outputs were stored on disk
- if the outputs fit in memory,
Why not store them in memory initially, only spill them to disk if they are large.
- Save data loading and storing time
- Shark's shuffle implementation does this by default

Data Format and Layout:

- MapReduce: pure schema-on-read approach
- Hive: itself supports "table partitions"
- Shark: using fast in-memory columnar representations within Spark.

Related Work

1. Compile declarative queries into MapReduce style jobs systems

Like ASTERIX, Tenzing, SCOPE, Cheetah and Hive

2. Implement low-latency engines using architectures resembling shared-nothing parallel databases

Like PowerDrill and Impala, Google's Dremel

3. hybrid approach by combining a MapReduce-like engine with relational databases

HadoopDB and Osprey

Answering

Question: What's the influence of Shark in the development of modern database?

Answer:

- (1) Shows that MapReduce-like execution models can be applied effectively to SQL
- (2) Offer a promising way to combine relational and complex analytics.

Question: What did Shark give us?

Answer:

- (1) Shark realize a low-latency system that can efficiently combine SQL and machine learning workloads.
- (2) Shark supports fine-grained fault recovery.

Question: Why Shark has been abandoned ?

Answer:

- (1) Shark inherited a large, complicated code base from Hive.
- (2) It made it hard to optimize and maintain. (constrained by Map-Reduce)
- (3) Ending development in Shark and move everything to Spark SQL.

Questions and Comments?

Thank you!